



COMP 478
CSUCI - Spring 2022

Lab7: A mini-project

Course Evaluation Result

- Thank You for your feedback!
- Additional resources
- Going through the lab's solution together in the class
- Not having HW and Lab in the same week
- More explanations on the lab assignments
- More examples
- Review session before the midterm

towards
data science

2 bonus points + midterm exam

Announcements

- The solution to Lab 5 is posted.
- The sample questions for the midterm exam are posted. I'll go through the solutions in our review session (Wed, March 23).
- Midterm Exam:
 - Lecture 1 – Lecture 13 (from the beginning to the end of the LR)

Let's do a mini project together!!

- Let's work on the “Wine recognition dataset” from sklearn!

A Real-world ML Problem

- Stroke Prediction Dataset:

- Binary classification task: +/-

- columns:

- 1) id: unique identifier
- 2) gender: "Male", "Female" or "Other"
- 3) age: age of the patient
- 4) hypertension: 0 if the patient doesn't have hypertension, 1 if the patient has hypertension
- 5) heart_disease: 0 if the patient doesn't have any heart diseases, 1 if the patient has a heart disease
- 6) ever_married: "No" or "Yes"
- 7) work_type: "children", "Govt_jov", "Never_worked", "Private" or "Self-employed"
- 8) Residence_type: "Rural" or "Urban"
- 9) avg_glucose_level: average glucose level in blood
- 10) bmi: body mass index
- 11) smoking_status: "formerly smoked", "never smoked", "smokes" or "Unknown"*
- 12) stroke: 1 if the patient had a stroke or 0 if not



Data Preprocessing

■ Load data:

```
import pandas as pd
```

```
df = pd.read_csv('./healthcare-dataset-stroke-data.csv')
```

healthcare-dataset-stroke-data

| id | gender | age | hypertension | heart_disease | ever_married | work_type | Residence_type | avg_glucose_level | bmi | smoking_status | stroke |
|-------|--------|-----|--------------|---------------|--------------|---------------|----------------|-------------------|------|-----------------|--------|
| 9046 | Male | 67 | 0 | 1 | Yes | Private | Urban | 228.69 | 36.6 | formerly smoked | 1 |
| 51676 | Female | 61 | 0 | 0 | Yes | Self-employed | Rural | 202.21 | N/A | never smoked | 1 |
| 31112 | Male | 80 | 0 | 1 | Yes | Private | Rural | 105.92 | 32.5 | never smoked | 1 |
| 60182 | Female | 49 | 0 | 0 | Yes | Private | Urban | 171.23 | 34.4 | smokes | 1 |
| 1665 | Female | 79 | 1 | 0 | Yes | Self-employed | Rural | 174.12 | 24 | never smoked | 1 |
| 56669 | Male | 81 | 0 | 0 | Yes | Private | Urban | 186.21 | 29 | formerly smoked | 1 |
| 53882 | Male | 74 | 1 | 1 | Yes | Private | Rural | 70.09 | 27.4 | never smoked | 1 |
| 10434 | Female | 69 | 0 | 0 | No | Private | Urban | 94.39 | 22.8 | never smoked | 1 |
| 27419 | Female | 59 | 0 | 0 | Yes | Private | Rural | 76.15 | N/A | Unknown | 1 |
| 60491 | Female | 78 | 0 | 0 | Yes | Private | Urban | 58.57 | 24.2 | Unknown | 1 |
| 12109 | Female | 81 | 1 | 0 | Yes | Private | Rural | 80.43 | 29.7 | never smoked | 1 |
| 12095 | Female | 61 | 0 | 1 | Yes | Govt_job | Rural | 120.46 | 36.8 | smokes | 1 |
| 12175 | Female | 54 | 0 | 0 | Yes | Private | Urban | 104.51 | 27.3 | smokes | 1 |
| 8213 | Male | 78 | 0 | 1 | Yes | Private | Urban | 219.84 | N/A | Unknown | 1 |
| 5317 | Female | 79 | 0 | 1 | Yes | Private | Urban | 214.09 | 28.2 | never smoked | 1 |
| 58202 | Female | 50 | 1 | 0 | Yes | Self-employed | Rural | 167.41 | 30.9 | never smoked | 1 |
| 56112 | Male | 64 | 0 | 1 | Yes | Private | Urban | 191.61 | 37.5 | smokes | 1 |

■ Input features and labels:

```
x = df[['gender', 'age', 'hypertension', 'heart_disease', 'ever_married',  
        'work_type', 'Residence_type', 'avg_glucose_level', 'bmi', 'smoking_status', 'stroke']]
```

```
print (x['age'].mean(), x['stroke'].value_counts())
```

Data Preprocessing

- Missing values:

1. Remove datapoints with missing values
2. Recover the value:
 - Continuous features: Average/mean, median, ...
 - Categorical features: Most frequent category
 - Predict missing values using classification, regression, ...
 -

```
df = df.dropna()
```

Data Preprocessing

- Handling non-numerical features:

```
x = pd.get_dummies(x)
print (x.columns)
```

```
Index(['age', 'hypertension', 'heart_disease', 'avg_glucose_level', 'bmi',
      'stroke', 'gender_Female', 'gender_Male', 'gender_Other',
      'ever_married_No', 'ever_married_Yes', 'work_type_Govt_job',
      'work_type_Never_worked', 'work_type_Private',
      'work_type_Self-employed', 'work_type_children', 'Residence_type_Rural',
      'Residence_type_Urban', 'smoking_status_Unknown',
      'smoking_status_formerly smoked', 'smoking_status_never smoked',
      'smoking_status_smokes'],
      dtype='object')
```

```
x = x.drop(columns='ever_married_Yes')
print (x.columns)
```

```
Index(['age', 'hypertension', 'heart_disease', 'avg_glucose_level', 'bmi',
      'stroke', 'gender_Female', 'gender_Male', 'gender_Other',
      'ever_married_No', 'work_type_Govt_job', 'work_type_Never_worked',
      'work_type_Private', 'work_type_Self-employed', 'work_type_children',
      'Residence_type_Rural', 'Residence_type_Urban',
      'smoking_status_Unknown', 'smoking_status_formerly smoked',
      'smoking_status_never smoked', 'smoking_status_smokes'],
      dtype='object')
```


Data Preprocessing

- Balance data: (Class +: 209, Class -: 4700)

How to balance data?

1. Down-sample the majority class

```
from sklearn.utils import resample

x_minority = x[x.stroke == 1]
x_majority = x[x.stroke == 0]

x_majority_downsampled = resample(x_majority, replace = False, n_samples = len(x_minority), random_state = 0)
```

2. Up-sample the minority class

```
from sklearn.utils import resample

x_minority = x[x.stroke == 1]
x_majority = x[x.stroke == 0]

x_minority_upsampled = resample(x_minority, replace = True, n_samples = len(x_majority) - len(x_minority), random_state = 0)
```

Training Phase

- Split data into train (70%) & test (30%):

```
Y = x_new['stroke']  
X = x_new.drop(columns='stroke')  
  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=0)
```

- KNN:

- The best model?
- The best result?

```
from sklearn.neighbors import KNeighborsClassifier  
neigh = KNeighborsClassifier(n_neighbors=1)  
neigh.fit(X_train, y_train)  
pred = neigh.predict(X_test)  
  
from sklearn.metrics import classification_report  
  
print (classification_report(y_test, pred))
```

How to Tune Threshold?

Predict → predicted class label

Predict_proba → probability associated to each class

| | pred |
|----|------|
| D1 | 1 |
| D2 | 1 |
| D3 | 0 |
| D4 | 1 |
| D5 | 0 |
| D6 | 0 |
| D7 | 0 |

| GT label |
|----------|
| 1 |
| 1 |
| 1 |
| 1 |
| 1 |
| 0 |
| 0 |

| | Class 0 | Class 1 |
|----|---------|---------|
| D1 | 0.09 | 0.91 |
| D2 | 0.36 | 0.64 |
| D3 | 0.55 | 0.45 |
| D4 | 0.3 | 0.7 |
| D5 | 0.59 | 0.41 |
| D6 | 0.95 | 0.05 |
| D7 | 0.87 | 0.13 |

ROC Curve

- Report the performance of a model:
 - Precision, Recall, f1score, accuracy, confusion matrix, ...
 - Receiver Operating Characteristic Curve (ROC curve):
 - Binary classification tasks
 - A graphical plot (TPR vs FPR) for different threshold
 - $TPR = \frac{TP}{TP+FN}$, $FPR = \frac{FP}{FP+TN}$

How to compare different models?
AUC (Area under the ROC Curve)

